

BidSense - Presentation Notes (speaker's talk track + Q&A)

Companion to gamma-outline.md (the 6 slides). This is what to SAY and how to handle the room. Built on the sales-playbook Value Equation and the label-then-answer objection structure. Spoken lines are judge-facing, so no em dashes, and every claim ties to something we actually built.

The core frame (read this first)

The whole pitch is one equation. Every sentence should raise the top or lower the bottom:

> **Value = (Dream Outcome x Perceived Likelihood) / (Time Delay x Effort)**

For BidSense, mapped:

Lever	Our line
Dream Outcome	Win more bids and never get disqualified on a technicality.
Perceived Likelihood	Trained on real data, validated in the open, every claim cited. It is proof, not promises.
Time Delay	Days of manual work become minutes. The demo proves it on the spot.
Effort	Upload one file. The analyst reviews and approves, the machine does the reading.

Positioning in one line: "Most AI proposal tools are a chatbot over a text box. BidSense is an auditable decision system that tells you GO or NO-GO before you waste a week."

The one thing they must remember: *It says no when it should.* If they forget everything else, they remember that we built the rare AI tool with the integrity to decline a bad bid.

The 30-second version (table visit, or if time gets cut)

"Bid teams lose 60 to 80 percent of their time just reading tenders, and one missed mandatory clause disqualifies the whole bid. BidSense reads a tender the way a senior bid manager does: it extracts every requirement, checks each one against your track record with cited evidence, scores your win probability with a real trained model, and drafts the proposal. You get a GO, CONDITIONAL GO, or NO-GO in minutes. And it is honest: it declined a road tender with four genuine gaps instead of flattering the bid. Want to see it run?"

Then open the demo.

The 5-minute talk track (slide by slide)

Timing target in brackets. Deliver outcome-first: lead with the number, then the mechanism.

Slide 1 - The problem [0:00 to 0:50]

"Responding to a public tender is slow, manual, and risky. A 15 to 80 page RFP takes two to four working days to reach a first draft. The mandatory requirements are scattered across eligibility clauses, submission rules, and annexes, and humans miss them. Miss one, and the entire bid is thrown out after days of work. Worse, nobody can tell you, before you spend that week, whether you can actually win. So firms chase bids they cannot win and skip ones they could."

Move: amplify the Dream Outcome by quantifying the leak. Do not pitch yet. Let the pain sit.

Slide 2 - The solution plus live demo [0:50 to 2:10]

"BidSense reads a tender the way a senior bid manager does. You upload a PDF or DOCX, and a four-stage pipeline runs: Extract every requirement, Match each one to your capability library with cited evidence, Score win probability, and Draft a compliant proposal. The output is one clear decision, GO, CONDITIONAL GO, or NO-GO, plus a review-ready draft. Let me show you."

Trigger the demo here. While it runs, narrate the stepper (see demo script below).

Move: compress Time Delay live. The audience watches days collapse into seconds.

Slide 3 - It says no when it should [2:10 to 3:00]

"Here is what makes this different. A tool that flatters every bid is useless. We ran a road tender with four genuine eligibility gaps, and BidSense said NO-GO and told us exactly why. Three things make that trustworthy. One, the win probability comes from a trained model with SHAP explanations, not a number a language model guessed. Two, every sentence in the draft cites a real capability record, and fabricated citations are rejected and unit-tested, so it cannot hallucinate a project you never did. Three, we publish our own ablation study inside the app instead of hiding the data's limits."

Move: maximize Perceived Likelihood. This is the integrity beat. Slow down here.

Slide 4 - Real ML, validated in the open [3:00 to 3:50]

"We trained on the hackathon's own data, a 120-bid history and a 50-record capability library. Every metric you see is out-of-fold, cross-validated, so it reflects generalization, not memorization. There is a real data-science decision here too: the awarded evaluation score is the only strong signal, but you only know it after evaluation, so it would be leakage to use it. We estimate it before submission from compliance, sector priors, and budget fit, then map that to win probability. It is all on a live validation page in the app, with the test suite. And it runs offline from a warmed cache, so it survives venue Wi-Fi."

Move: this is the proof slide for the technical judges. Point at the validation page if it is up.

Slide 5 - The market [3:50 to 4:30]

"This is a real market. Public procurement is 15 to 20 percent of GDP, and proposal-response software is already a global category with players like Responsive and Loopio. But none of them handle local reality: PPRA formats, PKR amounts, PEC categories, ISO requirements. Our beachhead is Pakistani IT-services and construction firms that bid on government and donor tenders every month. Cost per processed bid is a few cents of model spend, so the margins are software-grade. The model is per-seat SaaS plus per-bid pricing."

Move: lead with the size, land on the wedge no incumbent serves.

Slide 6 - Team, roadmap, and the ask [4:30 to 5:00]

"We are three: Aleem on data and ML, Anas on backend, Umer on frontend, and we built the full system end to end. Today it is a working single-company tool across three verticals. Next quarter, a

capability-library builder that ingests past proposals, a PPRA tender feed, and a win-loss loop that retrains on every closed bid. Year one, multi-tenant SaaS. Our ask is [INSERT YOUR ASK: pilot partners / mentorship / the prize]. BidSense turns a tender into a cited draft and an honest GO or NO-GO decision in minutes. Thank you."

Move: close on credibility plus momentum, then state the ask and stop talking.

The live demo script (the part that wins or loses the room)

Use a doc you have NOT pre-run if you want a fresh result, or the cached road tender for the guaranteed NO-GO reveal. Choreography:

1. **Upload.** "I am dropping in a real tender now." Drag the PDF onto the dropzone. 2. **Narrate the stepper while it runs.** "It is extracting every requirement, then matching each against our track record, then scoring, then drafting. Normally this is a two-day job." 3. **Land on the decision.** "There it is. [GO / NO-GO] at [X] percent win probability." Let the gauge sit on screen for a beat. 4. **The reveal (if NO-GO).** "Watch this. It declined the bid, and here are the four mandatory gaps it found. It is not guessing, each one points to a missing certification or a turnover threshold we do not meet. That honesty is the product." 5. **Open one drafted section.** "Every claim here is cited to a real capability record. No invented projects." 6. **Flash the validation page.** "And if you do not trust the model, here is the live proof: out-of-fold metrics, the ablation, the passing test suite."

Failsafe line if anything breaks or Wi-Fi dies: "Good news, this runs entirely offline from a warmed cache, so let me replay it from local." Then rerun. The resilience is itself a selling point, so a network failure becomes a feature demo, not a disaster.

Judge Q&A bank (label, then answer, then stop)

For each: acknowledge the question honestly first (never get defensive), give a tight answer, then stop talking. These are the questions the rubric and a sharp judge will actually ask.

"100 percent accuracy, isn't that overfitting?" Fair concern, it is the first thing we checked. The metric is out-of-fold: every bid is scored by a model that never trained on it, so it reflects generalization. It is near-perfect only because the awarded score nearly determines the outcome, which our ablation shows openly. A live bid's score is estimated, so real predictions carry that uncertainty and are not 100 percent confident.

"You synthesized the evaluation-criteria taxonomy. Isn't that made up?" Yes, and we disclose it on the slide and in the app. The sample dataset omitted the sheet the problem statement referenced, so we built a 16-entry taxonomy to fill the gap rather than pretend. Building it transparently is the honest move, and arguably a differentiator.

"Using the evaluation score sounds like circular reasoning." That is exactly the trap we avoided. The awarded score is only known after evaluation, so using it directly would be leakage. We never feed it live. We estimate the expected score before submission from levers the bidder controls, compliance, sector priors, budget fit, and the model maps that estimate to probability.

"n equals 120 is a tiny dataset." Agreed, and we treat it that way. We compare XGBoost against a simple regularized linear model on identical folds and ship whichever generalizes better, we report everything out-of-fold to avoid optimism, and we say so plainly on the validation page. The rigor is the

point.

"Why not just put GPT in a text box?" Because a chatbot will happily invent a project you never did and give you a confident number with no basis. We built three guardrails a chatbot does not have: a trained, explainable model for the probability, citations on every drafted claim with fabrication unit-tested, and a decision engine that says NO. That is auditable, which procurement actually requires.

"What stops Loopio or Responsive from doing this?" Two things. Localization they do not serve: PPRA, PKR, PEC categories, local ISO rules. And a data flywheel: as a firm processes bids, its capability library and win-loss history compound, which makes our scoring sharper for them specifically over time.

"Who pays, and how much?" Procurement and bid teams at firms that tender monthly. Per-seat SaaS plus per-bid pricing. The frame is simple: one disqualified bid costs days of senior time and a lost contract, so an annual fee that prevents even one pays for itself.

"Is this actually working or is it hardcoded for the demo?" Live, on data you can inspect. Upload a document we have not seen and it runs the full pipeline. The model, the tests, and the metrics are all in the repo and on the validation page.

"Can three people build this into a company?" We built the entire system, backend, ML, and frontend, inside the hackathon window, with clear ownership. The roadmap is sequenced so each step ships value on its own.

Delivery mechanics

- **Pace:** one idea per breath. Pause after the NO-GO reveal and after "it says no." Silence sells more than speed.
- **Outcome-first:** start sentences with the number or the result, then the mechanism. "Days become minutes, because the pipeline reads the whole tender for you."
- **Do not feature-dump.** Pick the three differentiators (explainable model, cited drafting, it says no) and repeat them. Drop the rest unless asked.
- **Hands off the keyboard during questions.** Answer, then stop. Do not fill silence with qualifications that weaken the answer.
- **The close is the ask plus silence.** State what you want, then let them respond.

Hard bans (will weaken the pitch)

- No "we think" or "we tried to" about things that work. Say what it does.
- No hedging on the honesty story. "It says no" is a strength, deliver it as one.
- No defensiveness in Q&A. Concede the real limits (n=120, synthesized taxonomy) plainly, that candor is what makes the rest credible.
- No reading the slides aloud. The slides are the backdrop, you are the signal.